

Coding Tasks on dplyr, tidyr and ggplot2 - 17th May 2021

Mario Keller

May 10, 2021

Contents

1	Introduction	2
2	Required libraries	2
3	dplyr	2
3.1	Task1 - Level: Beginner	2
3.2	Task2 - Level: Advanced	2
4	tidyr	3
4.1	Task3 - Level: Beginner	3
4.2	Task4 - Level: Advanced	3
4.3	Task5 - Level: Advanced	3
5	ggplot2	4
5.1	Task6 - Level: Beginner	4
5.2	Task7 - Level: Beginner	5
5.3	Task8 - Level: Advanced	5
5.4	Task9 - Level: Advanced	5
5.5	Task10 - Level: Advanced	6

1 Introduction

A good bioinformatician not only produces huge amounts of data, but is also able to cleanse and adjust these in order to then produce plots that are ready for publication. In today's coding task we will have a look on 3 R packages that are useful for the mentioned tasks.

The first of these packages is dplyr, which offers a grammar for data manipulation. The second package is tidyr, which helps you to clean up your data. The third package is ggplot2, which helps you to create good looking graphics that can be adjusted down to the smallest detail.

We will not fully explore the complete functionality of these packages but have a look on the most important functions. For further details check out the websites of the packages that are placed as a hyperlink in the package names above.

2 Required libraries

The following libraries and data are required in today's coding tasks.

```
library(dplyr)
library(tidyr)
library(ggplot2)
library(gridExtra)

#Adjust the path for your system
TPM.df <- readRDS("data/TPM.df.rds")
```

3 dplyr

One of the most important aspects about the dplyr packages is the so-called pipe operator `%>%`. This operator allows you to chain together sets of operations in a very intuitive fashion. Here is one example where we create a data frame, for which we calculate the mean for each row and afterwards the mean of the computed row means:

```
#With nested parentheses.
mean(rowMeans(data.frame(x=1:10, y=1:10)))

## [1] 5.5

#With %>%
data.frame(x=1:10, y=1:10) %>% rowMeans %>% mean

## [1] 5.5
```

As you can see, using the `%>%` operator is much more intuitive than using nested parentheses.

In addition the dplyr package offers additional functions that help you to filter your data, add new columns, adjust existing columns and arrange your data frames.

3.1 Task1 - Level: Beginner

Use the function `filter()` to filter for patients, whose Gene1 TPM values in brain and liver are > 150 and use afterwards `arrange()` to order the patients in descending order by their Gene1 Brain TPM value. Use the `%>%` operator and continue with the filtered data frame.

Hint: You should end up with 731 patients.

3.2 Task2 - Level: Advanced

Create a new column called `sumVar` which contains the sum of `Gene1_Brain`, `Gene2_Brain`, `Gene1_Liver` and `Gene2_Liver` for each patient. Afterwards filter for patients where `sumVar > 500` and subsequently remove the `sumVar` column with `select()`.

Helpful functions are `mutate()`, `rowwise()` and `c_across()`. Use the `%>%` operator.

Hint: Check out the example in `?c_across`. You should end up with 679 patients. The data frame is now a tibble, which is just an advanced data frame.

4 tidy

If you did not solve Task2 load the data.frame stored in “TPM.df.tidyr.rds” and continue with this data frame.

The tidy package offers different functions for cleaning up your data frames such as changing them from a wide to long format and vice versa.

4.1 Task3 - Level: Beginner

We will now mimic a scenario where for one of the patients the `Gene1_Brain` value is missing. Enter for the patient in row 3 an NA in the `Gene1_Brain` column (e.g. `TPM.df[3,2] <- NA` or `TPM.df$Gene1_Brain[3] <- NA`). Use `replace_na()` to replace the NA in the `Gene1_Brain` column with 361.0915.

Hint: `replace()` needs a list that contains the replacement values for each column.

4.2 Task4 - Level: Advanced

At the moment we have data frame in a wide format. However, we want it to be in a long format. Use `pivot_longer()` to produce a data frame where each combination of patient, gene and tissue is a single row.

Hint: You can split the names of the columns 2 to 5 by “_” and use “Gene” and “Tissue” as new column names. In addition the values column could be called “TPM”.

This is how the data frame should look like:

```
## # A tibble: 8 x 4
##   Patien.ID Gene  Tissue  TPM
##   <fct>      <chr> <chr>  <dbl>
## 1 164      Gene1 Brain   394.
## 2 164      Gene2 Brain    88.1
## 3 164      Gene1 Liver   189.
## 4 164      Gene2 Liver    82.4
## 5 747      Gene1 Brain   362.
## 6 747      Gene2 Brain    71.2
## 7 747      Gene1 Liver   195.
## 8 747      Gene2 Liver    91.9
```

4.3 Task5 - Level: Advanced

We can also go from a long format to a wide format. Use `pivot_wider()` to make two columns out of the “Gene” column.

Hint: The parameters “names_from” and “values_from” are helpful.

This is how the data frame should look like:

```
## # A tibble: 6 x 4
##   Patien.ID Tissue Gene1 Gene2
##   <fct>      <chr>  <dbl> <dbl>
## 1 164      Brain   394.  88.1
## 2 164      Liver   189.  82.4
## 3 747      Brain   362.  71.2
## 4 747      Liver   195.  91.9
## 5 842      Brain   361.  73.7
## 6 842      Liver   170.  76.4
```

5 ggplot2

If you did not solve Task5 load the data.frame stored in “TPM.df.ggplot2.rds” and continue with this data frame.

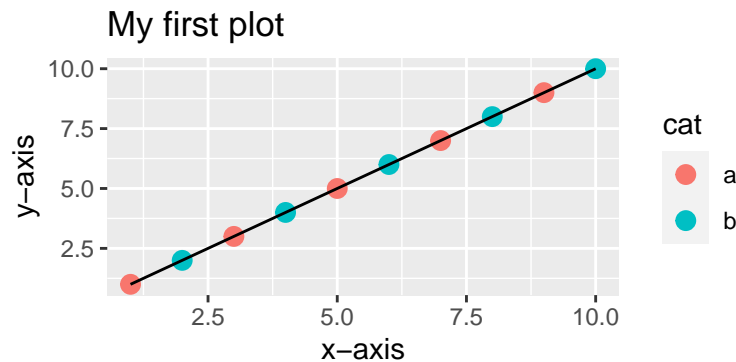
The ggplot2 package offers a modular structure that allows you to generate plots layer by layer. The most important function is `ggplot()` that initializes a ggplot object in which we declare the input data frame and set the plot aesthetics. The aesthetics are defined within `ggplot()` by using `aes()`. Aesthetics tell ggplot which columns contain the values on the x- and y-axis or which column should be used as a color or fill. After we have initialized a ggplot object with `ggplot()` we can add layers with the plus sign (+). Here are some examples of layers:

- `geom_point()` for scatter plots
- `geom_line()` for line plots
- `geom_col()` for bar charts
- `geom_histogram()` for histograms
- `geom_density()` for density plots

In addition there are different functions that can be used to adjust the x- and y-axis, to set labels or to set a user-defined color-scale.

Here is a simple example that shows how to create a scatter plot and to color the points by a certain column in the data.frame

```
#Create the data frame
data.frame(column1=1:10, column2=1:10, cat=factor(rep(c("a","b"),5), levels=c("a", "b")))) %>%
  #Initialize the ggplot object
  ggplot(., aes(x=column1, y=column2, color=cat)) +
  #Add a layer with points
  geom_point(size=3) +
  #Add a layer with a line
  geom_line(col="black") +
  #Re-name the axes and set a title
  labs(x="x-axis", y="y-axis", title="My first plot")
```



Once you have internalized the ggplot philosophy it is very easy to create beautiful plots that nicely represent your results.

5.1 Task6 - Level: Beginner

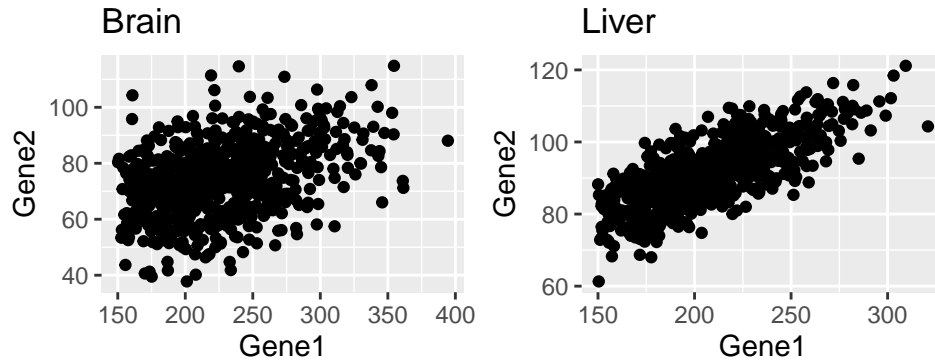
We want to make for each of the two tissues a scatter plot to show the correlation between Gene1 and Gene2. One easy way is to first subset the data.frame to have a single data frame for each of the tissues, create the plot and store it in a variable. Subset the data.frame once for the Brain tissue and once for the Liver tissue and store the two data frames as Brain.TPM.df and Liver.TPM.df

Hint: The `subset()` function is very useful.

5.2 Task7 - Level: Beginner


Use the two data frames to create two scatter plots using `geom_point()` with Gene1 on the x- and Gene 2 on the y-axis and store them as `Brain.plot` and `Liver.plot`. Afterwards use `grid.arrange()` from the `gridExtra` package to combine the two plots into a single plot side-by-side.

Hint: `x` and `y` are defined in `aes()`. You can tell `grid.arrange` whether you want the plots side by side or on top of each other (check out `ncol` and `nrow`).



5.3 Task8 - Level: Advanced


An alternative to the above approach is to use `facet_wrap()` from the `ggplot2` package to automatically create subplots. Use the `TPM.df` with `facet_wrap()` to create two subplots based on the `Tissue` column. In addition, color the the points by `Tissue`.



5.4 Task9 - Level: Advanced

As you can see the two plots share the same x- and y-axis. Try to set for each plot individual axes.

Hint: Check out `?facet_wrap` to find out how you can create for each plot individual x- and y-axes.



5.5 Task10 - Level: Advanced

When we try to look at correlations a regression line is often helpful. Use `geom_smooth()` to add a black regression line as a new layer.

Hint: Check out `?geom_smooth` and look for the correct “method”. `lm` could be useful.

